

Spark Basics



Agenda

- Hello World in Spark
- RDDs overview
- RDDs operations
-



Simple Spark Example

```
val conf = new SparkConf()  
  .setAppName("standalone app")
```

```
val sc = new SparkContext(conf)
```

```
val count = sc  
  .textFile("README.md")  
  .filter(l => l.contains("Python"))  
  .count()
```

```
println("Lines containing python: " + count)
```



Spark Basics - RDDs

- RDDs - Resilient Distributed Dataset
 - Immutable collection of objects, distributed across the cluster
 - Split into partitions to support parallelism
 - Create from external sources (HDFS, S3, etc) or in code
- RDD transformations
 - Return a new RDD (RDDs are immutable)
 - Lazily executed
 - Example: map, flatMap, filter
- RDD actions
 - Output to external systems (HDFS), or return to driver program
 - Force execution of (lazy) transformations
- Spark code just **builds** lineage graph (DAG of RDDs)
- RDDs can be recomputed and persisted



Spark Basics - RDDs (cont)

- Transformations

- Basic

- map - could return a different type
 - flatMap - fn returning iterator, "merges" the iterators across the elements
 - filter
 - sample(withRepl, frac, seed)

- Set-like

- distinct - slow, requires shuffle
 - union - fast, doesn't remove dupes
 - intersection - slow, also removes dupes (like distinct)
 - subtract

- cartesian - slow!



Spark Basics - RDDs (cont)

- Actions

- count : return total number elements in RDD
- collect : just take all of it and return to "driver"
- take(n) : minimizes num partitions it accesses => could have bias
- takeSample : like take but randomized
- top(n) : like take, but sorts first
- topOrdered(n, ordering) : custom ordering fn
- reduce : fold without initial
- fold : initial value + op
- aggregate : fold with combiner
- foreach : just run code, without data going to driver



Spark - Caching

- Example:

```
val result = input.map(x => x * x)
// result.persist(StorageLevel.DISK_ONLY)
println(result.count())
println(result.collect().mkString(", "))
```

- Levels

- MEMORY_ONLY
 - MEMORY_ONLY_SER
 - MEMORY_AND_DISK (spill to disk if mem is not enough, does not Serialize in memory)
 - MEMORY_AND_DISK_SER (spill to disk if mem is not enough, does Serialize in memory)
 - DISK_ONLY
- on eviction, it uses LRU, del or disk
 - unpersist -- manually evict something from cache



Spark Basics - RDDs

- Pair transformations (**PairRDDFunctions**)
 - reduceByKey : like reduce, but acts per key
 - foldByKey : like fold, but acts per key
 - combineByKey(createCombiner, mergeValue, mergeCombiners) : like aggregate
 - parallelism/partitions can be configured on all combine* transformations
 - mapValues : preserves partition
 - flatMapValues
 - keys
 - values
 - sortByKey
 - join, leftOuterJoin, rightOuterJoin, cogroup
 - repartition
 - Coalesce : like repartition but only reduces





Spark Basics - Pair RDD partitioning

- Defines parallelism
- Defines data locality
- Crucial for performance optimizations
- Use `persist()` after `partitionBy`, otherwise it will be recomputed
- `sortByKey` results in range-partitioned RDD
- `groupByKey` results in hash-partitioned RDD
- `map()` creates RDD without partition info (i.e “erases” it)



Spark Basics - Partitioning example



Spark Basics - Partitioning details

- benefit from partitioning:
 - cogroup, groupWith, join*, groupByKey, combineByKey*, lookup
- will set a partitioner:
 - cogroup, groupWith, join*, groupByKey, combineByKey*, partitionBy, sort
 - mapValues, flatMapValues, filter (only if parent had partition)
- binary partitioners will use the partition of one of their operands, or the first one, if both have partitioners
- PageRank example



Links

Repo with code and other lectures: <https://github.com/ichernev/spark-course>

